

Efficient Selective-ID Secure Identity Based Encryption Without Random Oracles

Dan Boneh*	Xavier Boyen
Stanford University	Voltage Security
<i>dabo@cs.stanford.edu</i>	<i>xb@boyen.org</i>

September 2004

Abstract

We construct two efficient Identity Based Encryption (IBE) systems that are selective identity secure *without the random oracle model* in groups equipped with a bilinear map. Selective identity secure IBE is a slightly weaker security model than the standard security model for IBE. In this model the adversary must commit ahead of time to the identity that it intends to attack, whereas in the standard model the adversary is allowed to choose this identity adaptively. The first system is based on the decisional bilinear Diffie-Hellman assumption, and extends to give a selective identity Hierarchical IBE secure without random oracles. The second system is based on a related assumption called the bilinear Diffie-Hellman inversion assumption. Applications of either system include an efficient CCA2 public key cryptosystem that supports non-interactive threshold decryption in the standard model, and a simple and practical IBE system that remains secure against full adaptive-ID attacks, under some security penalty, without random oracles.

Keywords: identity-based encryption, bilinear groups, selective-ID security, full IBE security, CCA2 security, bilinear Diffie-Hellman inversion, public-key encryption, probable security.

1 Introduction

Identity Based Encryption (IBE) provides a public key encryption mechanism where a public key is an arbitrary string such as an email address or a telephone number. The corresponding private key can only be generated by a Private Key Generator (PKG) who has knowledge of a master secret. In an IBE system, users authenticate themselves to the PKG and obtain private keys corresponding to their identities. The identity based encryption model was first proposed without solution two decades ago [29], and several different approaches were subsequently suggested in a few precursor papers [30, 31, 20]. It is only a few years ago, however, that a formal security model and a practical implementation were proposed. Boneh and Franklin [5, 6] define a security model for identity based encryption and gave a construction based on the Bilinear Diffie-Hellman (BDH) problem. Cocks [12] describes another construction using quadratic residues modulo a composite, which however is less practical than the Boneh-Franklin system. The security of these systems requires cryptographic hash functions that are modeled as random oracles, i.e., these systems are only proven secure in the random oracle model [1]. A natural question is to devise a secure IBE system without random oracles.

*Supported by NSF and the Packard Foundation.

In the Boneh-Franklin security model the adversary can issue both adaptive chosen ciphertext queries and adaptive chosen identity queries (i.e., the adversary can request the private key for identities of its choice). Eventually, the adversary adaptively chooses the identity it wishes to attack and asks for a semantic security challenge for this identity. Canetti et al. [10, 11] recently proposed a slightly weaker security model, called selective identity IBE. In this model the adversary must commit ahead of time (non-adaptively) to the identity it intends to attack. The adversary can still issue adaptive chosen ciphertext and adaptive chosen identity queries. Canetti et al. are able to construct a provably secure IBE in this weaker model without the random oracle model. However, their construction views identities as bit strings, causing their system to require a bilinear map computation for every bit in the identity.

We construct two efficient IBE systems that are provably selective identity secure without the random oracle model. In both systems, encryption requires no bilinear map computation and decryption requires at most two. Our first construction is based on the Decision Bilinear Diffie-Hellman (Decision BDH) assumption. This construction extends to give an efficient selective identity secure Hierarchical IBE (HIBE) without random oracles. Hierarchical IBE was defined in [15] and the first construction in the random oracle model was given by Gentry and Silverberg [14]. Our efficient HIBE construction is similar to the Gentry-Silverberg system, but we are able to prove security without using random oracles. Our second IBE construction is even more efficient, but is based on a non-standard assumption we call Decision Bilinear Diffie-Hellman Inversion (Decision BDHI). Roughly speaking, the assumption says that no efficient algorithm can distinguish $e(g, g)^{1/x}$ from random, given $g, g^x, g^{(x^2)}, \dots, g^{(x^q)}$ for some q .

Canetti et al. [11] recently showed that any selective identity, chosen plaintext IBE gives a chosen ciphertext secure (CCA2) public key system. Consequently, both our IBE systems give efficient CCA2-secure public key systems without random oracles. Performance of both these CCA2-secure systems is comparable to the performance of the Cramer-Shoup system [13] which is based on Decision Diffie-Hellman.

Towards our goal of providing a practical secure IBE system without random oracle, we show that any selective-ID secure IBE can be turned into a fully secure IBE in the standard model, albeit at the cost of an inefficient security reduction. We note that Boneh and Boyen [4] recently gave the first fully secure IBE construction that is polynomial in all parameters, a construction which was subsequently simplified and made much more efficient by Waters [32]. For practical purposes, the transformation from efficient selective-ID IBE systems presented here results in even simpler and more efficient systems: indeed, the large security penalty incurred by the transformation is easily compensated for by increasing the size of the groups used by the system. The result are two very simple secure IBE systems without random oracles.

2 Preliminaries

Before presenting our results we briefly review the definition of security for an IBE system. We also review the definition of groups equipped with a bilinear map.

2.1 Selective Identity Secure IBE and HIBE Systems

Recall that an Identity Based Encryption system (IBE) consists of four algorithms [29, 5]: *Setup*, *KeyGen*, *Encrypt*, *Decrypt*. The *Setup* algorithm generates system parameters, denoted by *params*, and a master key *master-key*. The *KeyGen* algorithm uses the master key to generate the private key corresponding to a given identity. The encryption algorithm encrypts messages for a given

identity (using the system parameters) and the decryption algorithm decrypts ciphertexts using the private key. In a Hierarchical IBE [15, 14] identities are vectors. A vector of dimension ℓ represents an identity at depth ℓ . Algorithm *KeyGen* takes as input an identity $ID = (I_1, \dots, I_\ell)$ at depth ℓ and the private key $d_{ID|_{\ell-1}}$ of the parent identity $ID|_{\ell-1} = (I_1, \dots, I_{\ell-1})$ at depth $\ell - 1$. It outputs the private key d_{ID} for identity ID . We refer to the *master-key* as the private key at depth 0 and note that an IBE system is an HIBE where all identities are at depth 1.

Boneh and Franklin [5, 6] define chosen ciphertext security for IBE systems under a chosen identity attack. In their model the adversary is allowed to adaptively chose the public key it wishes to attack (the public key on which it will be challenged). Canetti, Halevi, and Katz [10, 11] define a weaker notion of security in which the adversary commits ahead of time to the public key it will attack. We refer to this notion as selective identity, chosen ciphertext secure IBE (IND-sID-CCA). More precisely, selective identity IBE and HIBE security is defined using the following game:

Init: The adversary outputs an identity ID^* where it wishes to be challenged.

Setup: The challenger runs the *Setup* algorithm. It gives the adversary the resulting system parameters *params*. It keeps the *master-key* to itself.

Phase 1: The adversary issues queries q_1, \dots, q_m where query q_i is one of:

- Private key query $\langle ID_i \rangle$ where $ID_i \neq ID^*$ and ID_i is not a prefix of ID^* . The challenger responds by running algorithm *KeyGen* to generate the private key d_i corresponding to the public key $\langle ID_i \rangle$. It sends d_i to the adversary.
- Decryption query $\langle C_i \rangle$ for identity ID^* or any prefix of ID^* . The challenger responds by running algorithm *KeyGen* to generate the private key d corresponding to ID^* (or the relevant prefix thereof as requested). It then runs algorithm *Decrypt* to decrypt the ciphertext C_i using the private key d . It sends the resulting plaintext to the adversary.

These queries may be asked adaptively, that is, each query q_i may depend on the replies to q_1, \dots, q_{i-1} .

Challenge: Once the adversary decides that Phase 1 is over it outputs two equal length plaintexts $M_0, M_1 \in \mathcal{M}$ on which it wishes to be challenged. The challenger picks a random bit $b \in \{0, 1\}$ and sets the challenge ciphertext to $C = \text{Encrypt}(\text{params}, ID^*, M_b)$. It sends C as the challenge to the adversary.

Phase 2: The adversary issues additional queries q_{m+1}, \dots, q_n where q_i is one of:

- Private key query $\langle ID_i \rangle$ where $ID_i \neq ID^*$ and ID_i is not a prefix of ID^* . The challenger responds as in Phase 1.
- Decryption query $\langle C_i \rangle \neq \langle C \rangle$ for ID^* or any prefix of ID^* . The challenger responds as in Phase 1.

These queries may be asked adaptively as in Phase 1.

Guess: Finally, the adversary outputs a guess $b' \in \{0, 1\}$. The adversary wins if $b = b'$.

We refer to such an adversary \mathcal{A} as an IND-sID-CCA adversary. We define the advantage of the adversary \mathcal{A} in attacking the scheme \mathcal{E} as

$$\text{Adv}_{\mathcal{E}, \mathcal{A}} = \left| \Pr[b = b'] - \frac{1}{2} \right|$$

The probability is over the random bits used by the challenger and the adversary.

Definition 2.1. We say that an IBE or HIBE system \mathcal{E} is $(t, q_{\text{ID}}, q_C, \epsilon)$ -selective identity, adaptive chosen ciphertext secure if for any t -time IND-sID-CCA adversary \mathcal{A} that makes at most q_{ID} chosen private key queries and at most q_C chosen decryption queries we have that $\text{Adv}_{\mathcal{E}, \mathcal{A}} < \epsilon$. As shorthand, we say that \mathcal{E} is $(t, q_{\text{ID}}, q_C, \epsilon)$ IND-sID-CCA secure.

Semantic Security. As usual, we define selective identity, chosen plaintext security for an IBE system as in the preceding game, except that the adversary is not allowed to issue any decryption queries. The adversary may still issue adaptive private key queries.

Definition 2.2. We say that an IBE or HIBE system \mathcal{E} is $(t, q_{\text{ID}}, \epsilon)$ -selective identity, chosen plaintext secure if \mathcal{E} is $(t, q_{\text{ID}}, 0, \epsilon)$ -selective identity, chosen ciphertext secure. As shorthand, we say that \mathcal{E} is $(t, q_{\text{ID}}, \epsilon)$ IND-sID-CPA secure.

2.2 Bilinear Groups

We briefly review the necessary facts about bilinear maps and bilinear map groups. We use the following notation:

1. \mathbb{G} and \mathbb{G}_1 are two (multiplicative) cyclic groups of prime order p ;
2. g is a generator of \mathbb{G} .
3. e is a bilinear map $e : \mathbb{G} \times \mathbb{G} \rightarrow \mathbb{G}_1$.

Let \mathbb{G} and \mathbb{G}_1 be two groups as above. A bilinear map is a map $e : \mathbb{G} \times \mathbb{G} \rightarrow \mathbb{G}_1$ with the following properties:

1. Bilinearity: for all $u, v \in \mathbb{G}$ and $a, b \in \mathbb{Z}$, we have $e(u^a, v^b) = e(u, v)^{ab}$.
2. Non-degeneracy: $e(g, g) \neq 1$.

We say that \mathbb{G} is a bilinear group if the group action in \mathbb{G} can be computed efficiently and there exists a group \mathbb{G}_1 and an efficiently computable bilinear map $e : \mathbb{G} \times \mathbb{G} \rightarrow \mathbb{G}_1$ as above. Note that $e(\cdot, \cdot)$ is symmetric since $e(g^a, g^b) = e(g, g)^{ab} = e(g^b, g^a)$.

Throughout the paper, for a prime order group \mathbb{G} we use \mathbb{G}^* to denote the set $\mathbb{G} \setminus \{1_{\mathbb{G}}\}$ where $1_{\mathbb{G}}$ is the identity of \mathbb{G} .

3 Complexity Assumptions

Let \mathbb{G} be a bilinear group of prime order p . We review the standard Bilinear Diffie-Hellman (BDH) assumption and define the Bilinear Diffie-Hellman Inversion (BDHI) assumption.

3.1 Bilinear Diffie-Hellman Assumption

The BDH problem [16, 28, 5] in \mathbb{G} is as follows: given a tuple $g, g^a, g^b, g^c \in \mathbb{G}$ as input, output $e(g, g)^{abc} \in \mathbb{G}_1$. An algorithm \mathcal{A} has advantage ϵ in solving BDH in \mathbb{G} if

$$\Pr \left[\mathcal{A}(g, g^a, g^b, g^c) = e(g, g)^{abc} \right] \geq \epsilon$$

where the probability is over the random choice of generator g in \mathbb{G}^* , the random choice of a, b, c in \mathbb{Z}_p , and the random bits used by \mathcal{A} . Similarly, we say that an algorithm \mathcal{B} that outputs $b \in \{0, 1\}$ has advantage ϵ in solving the *decision* BDH problem in \mathbb{G} if

$$\left| \Pr \left[\mathcal{B}(g, g^a, g^b, g^c, e(g, g)^{abc}) = 0 \right] - \Pr \left[\mathcal{B}(g, g^a, g^b, g^c, T) = 0 \right] \right| \geq \epsilon$$

where the probability is over the random choice of generator g in \mathbb{G}^* , the random choice of a, b, c in \mathbb{Z}_p , the random choice of $T \in \mathbb{G}_1$, and the random bits consumed by \mathcal{B} . We refer to the distribution on the left as \mathcal{P}_{BDH} and the distribution on the right as \mathcal{R}_{BDH} .

Definition 3.1. We say that the (Decision) (t, ϵ) -BDH assumption holds in \mathbb{G} if no t -time algorithm has advantage at least ϵ in solving the (Decision) BDH problem in \mathbb{G} .

Occasionally we drop the t and ϵ and refer to the BDH and Decision BDH assumptions in \mathbb{G} .

3.2 Bilinear Diffie-Hellman Inversion Assumption

The q -BDHI problem is defined as follows: given the $(q+1)$ -tuple $(g, g^x, g^{(x^2)}, \dots, g^{(x^q)}) \in (\mathbb{G}^*)^{q+1}$ as input, compute $e(g, g)^{1/x} \in \mathbb{G}_1^*$. An algorithm \mathcal{A} has advantage ϵ in solving q -BDHI in \mathbb{G} if

$$\Pr \left[\mathcal{A}(g, g^x, \dots, g^{(x^q)}) = e(g, g)^{1/x} \right] \geq \epsilon$$

where the probability is over the random choice of generator g in \mathbb{G}^* , the random choice of x in \mathbb{Z}_p^* , and the random bits of \mathcal{A} . Similarly, we say that an algorithm \mathcal{B} that outputs $b \in \{0, 1\}$ has advantage ϵ in solving the *decision* q -BDHI problem in \mathbb{G} if

$$\left| \Pr \left[\mathcal{B}(g, g^x, \dots, g^{(x^q)}, e(g, g)^{1/x}) = 0 \right] - \Pr \left[\mathcal{B}(g, g^x, \dots, g^{(x^q)}, T) = 0 \right] \right| \geq \epsilon$$

where the probability is over the random choice of generator g in \mathbb{G}^* , the random choice of x in \mathbb{Z}_p^* , the random choice of $T \in \mathbb{G}_1$, and the random bits of \mathcal{B} . We refer to the distribution on the left as \mathcal{P}_{BDHI} and the distribution on the right as \mathcal{R}_{BDHI} .

Definition 3.2. We say that the (Decision) (t, q, ϵ) -BDHI assumption holds in \mathbb{G} if no t -time algorithm has advantage at least ϵ in solving the (Decision) q -BDHI problem in \mathbb{G} .

Occasionally we drop the t and ϵ and refer to the q -BDHI and Decision q -BDHI assumptions. It is easy to show that the 1-BDHI assumption is equivalent to the standard Bilinear Diffie-Hellman assumption (BDH). It is not known if the q -BDHI assumption, for $q > 1$, is equivalent to BDH. A closely related assumption was previously used in [22] where it was called weak Diffie-Hellman.

4 Efficient Selective Identity IBE and HIBE Based on BDH Without Random Oracles

We construct an efficient HIBE system that is selective identity secure without random oracles based on the Decision BDH assumption. In particular, this implies an efficient selective identity, chosen ciphertext secure IBE based on Decision BDH without random oracles.

4.1 Construction

Let \mathbb{G} be a bilinear group of prime order p (the security parameter determines the size of \mathbb{G}). Let $e : \mathbb{G} \times \mathbb{G} \rightarrow \mathbb{G}_1$ be the bilinear map. For now, we assume public keys (ID) of depth ℓ are vectors of elements in \mathbb{Z}_p^ℓ . We write $\text{ID} = (\text{I}_1, \dots, \text{I}_\ell) \in \mathbb{Z}_p^\ell$. The j -th component corresponds to the identity at level j . We later extend the construction to public keys over $\{0, 1\}^*$ by first hashing each component I_j using a collision resistant hash $H : \{0, 1\}^* \rightarrow \mathbb{Z}_p$. We also assume messages to be encrypted are elements in \mathbb{G}_1 . The HIBE system works as follows:

Setup(ℓ): To generate system parameters for an HIBE of maximum depth ℓ , select a random generator g in \mathbb{G}^* , a random $\alpha \in \mathbb{Z}_p$, and set $g_1 = g^\alpha$. Next, pick random elements $h_1, \dots, h_\ell \in \mathbb{G}$ and a random element $g_2 \in \mathbb{G}$. The public parameters *params* and the secret *master-key* are given by

$$params = (g, g_1, g_2, h_1, \dots, h_\ell), \quad master\text{-}key = g_2^\alpha$$

For $j = 1, \dots, \ell$, we define $F_j : \mathbb{Z}_p \rightarrow \mathbb{G}$ to be the function: $F_j(x) = g_1^x h_j$.

KeyGen($d_{ID|j-1}, ID$): To generate the private key d_{ID} for an identity $ID = (I_1, \dots, I_j) \in \mathbb{Z}_p^j$ of depth $j \leq \ell$, pick random $r_1, \dots, r_j \in \mathbb{Z}_p$ and output

$$d_{ID} = \left(g_2^\alpha \cdot \prod_{k=1}^j F_k(I_k)^{r_k}, \quad g^{r_1}, \quad \dots, \quad g^{r_j} \right)$$

Note that the private key for ID can be generated just given a private key for $ID_{|j-1} = (I_1, \dots, I_{j-1}) \in \mathbb{Z}_p^{j-1}$, as required. Indeed, let $d_{ID|j-1} = (d_0, \dots, d_{j-1})$ be the private key for $ID_{|j-1}$. To generate d_{ID} pick a random $r_j \in \mathbb{Z}_p$ and output $d_{ID} = (d_0 \cdot F_j(I_j)^{r_j}, d_1, \dots, d_{j-1}, g^{r_j})$.

Encrypt(*params*, ID , M): To encrypt a message $M \in \mathbb{G}_1$ under the public key $ID = (I_1, \dots, I_j) \in \mathbb{Z}_p^j$, pick a random $s \in \mathbb{Z}_p$ and output

$$C = \left(e(g_1, g_2)^s \cdot M, \quad g^s, \quad F_1(I_1)^s, \quad \dots, \quad F_j(I_j)^s \right)$$

Note that $e(g_1, g_2)$ can be precomputed once and for all so that encryption does not require any pairing computations. Alternatively, $e(g_1, g_2)$ can be included in the system parameters, in which case g_2 can be dropped.

Decrypt(d_{ID}, C): Consider an identity $ID = (I_1, \dots, I_j)$. To decrypt a given ciphertext $C = (A, B, C_1, \dots, C_j)$ using the private key $d_{ID} = (d_0, d_1, \dots, d_j)$, output

$$A \cdot \frac{\prod_{k=1}^j e(C_k, d_k)}{e(B, d_0)} = M$$

Indeed, for a valid ciphertext, we have

$$\frac{\prod_{k=1}^j e(C_k, d_k)}{e(B, d_0)} = \frac{\prod_{k=1}^j e(F_k(I_k), g)^{sr_k}}{e(g, g_2)^{s\alpha} \prod_{k=1}^j e(g, F_k(I_k))^{sr_k}} = \frac{1}{e(g_1, g_2)^s}$$

4.2 Security

The HIBE system above is reminiscent of the Gentry-Silverberg HIBE which is only known to be secure in the random oracle model. Surprisingly, our choice of functions F_1, \dots, F_ℓ enables us to prove security *without random oracles*. We prove security of our HIBE system under the standard Decision BDH assumption in \mathbb{G} .

Theorem 4.1. *Suppose the (t, ϵ) -Decision BDH assumption holds in \mathbb{G} . Then the previously defined ℓ -HIBE system is (t', q_s, ϵ) -selective identity, chosen plaintext (IND-sID-CPA) secure for arbitrary ℓ and q_s , and any $t' < t - \Theta(\tau \ell q_s)$ where τ is the maximum time for an exponentiation in \mathbb{G} .*

Proof. Suppose \mathcal{A} has advantage ϵ in attacking the HIBE system. We build an algorithm \mathcal{B} that solves the Decision BDH problem in \mathbb{G} . Algorithm \mathcal{B} is given as input a random 5-tuple (g, g^a, g^b, g^c, T) that is either sampled from \mathcal{P}_{BDH} (where $T = e(g, g)^{abc}$) or from \mathcal{R}_{BDH} (where T is uniform and independent in \mathbb{G}_1). Algorithm \mathcal{B} 's goal is to output 1 if $T = e(g, g)^{abc}$ and 0 otherwise. Set $g_1 = g^a$, $g_2 = g^b$, $g_3 = g^c$. Algorithm \mathcal{B} works by interacting with \mathcal{A} in a selective identity game as follows:

Initialization. The selective identity game begins with \mathcal{A} first outputting an identity $\text{ID}^* = (I_1^*, \dots, I_k^*) \in \mathbb{Z}_p^k$ of depth $k \leq \ell$ that it intends to attack. If necessary, \mathcal{B} appends random elements in \mathbb{Z}_p to ID^* so that ID^* is a vector of length ℓ .

Setup. To generate the system parameters, algorithm \mathcal{B} picks $\alpha_1, \dots, \alpha_\ell \in \mathbb{Z}_p$ at random and defines $h_j = g_1^{-I_j^*} g^{\alpha_j} \in \mathbb{G}$ for $j = 1, \dots, \ell$. It gives \mathcal{A} the system parameters $params = (g, g_1, g_2, h_1, \dots, h_\ell)$. Note that the corresponding master key, which is unknown to \mathcal{B} , is $g_2^a = g^{ab} \in \mathbb{G}$. As before, for $j = 1, \dots, \ell$ we define $F_j : \mathbb{Z}_p \rightarrow \mathbb{G}$ to be the function

$$F_j(x) = g_1^x h_j = g_1^{x-I_j^*} g^{\alpha_j}$$

Phase 1. \mathcal{A} issues up to q_S private key queries. Consider a query for the private key corresponding to $\text{ID} = (I_1, \dots, I_u) \in \mathbb{Z}_p^u$ where $u \leq \ell$. The only restriction is that ID is not a prefix of ID^* . Let j be the smallest index such that $I_j \neq I_j^*$. Necessarily $1 \leq j \leq u$. To respond to the query, algorithm \mathcal{B} first derives a private key for the identity (I_1, \dots, I_j) from which it then constructs a private key for the requested identity $\text{ID} = (I_1, \dots, I_j, \dots, I_u)$. Algorithm \mathcal{B} picks random elements $r_1, \dots, r_j \in \mathbb{Z}_p$ and sets

$$d_0 = g_2^{\frac{-\alpha_j}{I_j - I_j^*}} \prod_{v=1}^j F_v(I_v)^{r_v}, \quad d_1 = g^{r_1}, \quad \dots, \quad d_{j-1} = g^{r_{j-1}}, \quad d_j = g_2^{\frac{-1}{I_j - I_j^*}} g^{r_j}$$

We claim that (d_0, d_1, \dots, d_j) is a valid random private key for (I_1, \dots, I_j) . To see this, let $\tilde{r}_j = r_j - b/(I_j - I_j^*)$. Then we have that

$$g_2^{\frac{-\alpha_j}{(I_j - I_j^*)}} F_j(I_j)^{r_j} = g_2^{\frac{-\alpha_j}{(I_j - I_j^*)}} (g_1^{I_j - I_j^*} g^{\alpha_j})^{r_j} = g_2^a (g_1^{I_j - I_j^*} g^{\alpha_j})^{r_j - \frac{b}{I_j - I_j^*}} = g_2^a F_j(I_j)^{\tilde{r}_j}$$

It follows that the private key (d_0, d_1, \dots, d_j) defined above satisfies

$$d_0 = g_2^a \cdot \left(\prod_{v=1}^{j-1} F_v(I_v)^{r_v} \right) \cdot F_j(I_j)^{\tilde{r}_j}, \quad d_1 = g^{r_1}, \quad \dots, \quad d_{j-1} = g^{r_{j-1}}, \quad d_j = g^{\tilde{r}_j}$$

where $r_1, \dots, r_{j-1}, \tilde{r}_j$ are uniform in \mathbb{Z}_p . This matches the definition for a private key for (I_1, \dots, I_j) . Hence, (d_0, d_1, \dots, d_j) is a valid private key for (I_1, \dots, I_j) . Algorithm \mathcal{B} derives a private key for the requested ID from the private key (d_0, d_1, \dots, d_j) and gives \mathcal{A} the result.

Challenge. When \mathcal{A} decides that Phase 1 is over, it outputs two messages $M_0, M_1 \in \mathbb{G}_1$ on which it wishes to be challenged. Algorithm \mathcal{B} picks a random bit $b \in \{0, 1\}$ and responds with the ciphertext $C = (M_b \cdot T, g_3, g_3^{\alpha_1}, \dots, g_3^{\alpha_k})$. Since $F_i(I_i^*) = g^{\alpha_i}$ for all i , we have that

$$C = (M_b \cdot T, g^c, F_1(I_1^*)^c, \dots, F_k(I_k^*)^c)$$

Hence, if $T = e(g, g)^{abc} = e(g_1, g_2)^c$ then C is a valid encryption of M_b under the public key $\text{ID}^* = (I_1^*, \dots, I_k^*)$. On the other hand, when T is uniform and independent in \mathbb{G}_1 (when the input 5-tuple is sampled from \mathcal{R}_{BDH}) then C is independent of b in the adversary's view.

Phase 2. \mathcal{A} continues to issue queries not issued in Phase 1. Algorithm \mathcal{B} responds as before.

Guess. Finally, \mathcal{A} outputs a guess $b' \in \{0, 1\}$. Algorithm \mathcal{B} concludes its own game by outputting a guess as follows. If $b = b'$ then \mathcal{B} outputs 1 meaning $T = e(g, g)^{abc}$. Otherwise, it outputs 0 meaning $T \neq e(g, g)^{abc}$.

When the input 5-tuple is sampled from \mathcal{P}_{BDH} (where $T = e(g, g)^{abc}$) then \mathcal{A} 's view is identical to its view in a real attack game and therefore \mathcal{A} must satisfy $|\Pr[b = b'] - 1/2| > \epsilon$. On the other hand, when the input 5-tuple is sampled from \mathcal{R}_{BDH} (where T is uniform in \mathbb{G}_1) then $\Pr[b = b'] = 1/2$. Therefore, with g uniform in \mathbb{G}^* , a, b, c uniform in \mathbb{Z}_p , and T uniform in \mathbb{G}_1 we have that

$$\left| \Pr \left[\mathcal{B}(g, g^a, g^b, g^c, e(g, g)^{abc}) = 0 \right] - \Pr \left[\mathcal{B}(g, g^a, g^b, g^c, T) = 0 \right] \right| \geq \left| \left(\frac{1}{2} \pm \epsilon \right) - \frac{1}{2} \right| = \epsilon$$

as required. This completes the proof of Theorem 4.1. \square

4.3 Chosen Ciphertext Security

A recent result of Canetti et al. [11] gives an efficient way to build a selective identity, chosen ciphertext ℓ -HIBE from a selective identity, chosen plaintext $(\ell + 1)$ -HIBE. In combination with the above construction, we obtain a selective identity, chosen ciphertext ℓ -HIBE for any ℓ . In particular, from our 2-HIBE we obtain an efficient selective identity, chosen ciphertext secure IBE without random oracles.

4.4 Arbitrary Identities

We can extend our HIBE above to handle identities $ID = (I_1, \dots, I_\ell)$ with $I_j \in \{0, 1\}^*$ (as opposed to $I_j \in \mathbb{Z}_p$) by first hashing each I_j using a collision resistant hash function $H : \{0, 1\}^* \rightarrow \mathbb{Z}_p$ prior to key generation and encryption. A standard argument shows that if the scheme above is selective identity, chosen ciphertext secure then so is the scheme with the additional hash function. We note that there is no need for a full domain hash into \mathbb{Z}_p ; for example, a collision resistant hash function $H : \{0, 1\}^* \rightarrow \{1, \dots, 2^b\}$ where $2^b < p$ is sufficient for the security proof.

4.5 Fast Key Generation

It is worth noting that in the practical case of a non-hierarchical IBE system without secret sharing, the *KeyGen* algorithm of Section 4.1 can be implemented very efficiently. The implementation makes use of the discrete logarithms of g_1 and h_1 , which is accommodated by slightly modifying the *Setup* algorithm to select $h_1 \leftarrow g^\eta$ where $\eta \in \mathbb{Z}_p$ is random, and to preserve the (secret) values of α and η along with the master key g_2^α for future use. Under these assumptions, the expression of the private key d_{ID} returned by *KeyGen* for a given identity ID can be calculated as follows:

$$d_{ID} = \left(g_2^\alpha h_1^r g_1^{r \cdot ID}, g^r \right) = \left(g_2^\alpha \cdot g^{(\eta + \alpha \cdot ID)r}, g^r \right) = \left(\text{master-key} \cdot g^\xi, g^r \right)$$

Since $\text{master-key} = g_2^\alpha$ is constant, the only variable quantities in the last expression are the two powers g^ξ and g^r of the same fixed base g to the known exponents r and $\xi = (\eta + \alpha \cdot ID)r$ respectively. Without precomputation, the cost of jointly computing both quantities slightly exceeds that of a single exponentiation in \mathbb{G} , which for the sake of comparison is also the cost of key generation in

the Boneh-Franklin IBE system. Furthermore, since here the exponentiation base g is fixed, pre-computation techniques can be exploited to further reduce the total cost to a fraction of that level. For example, using an exponentiation algorithm with W -bit windows, the asymptotic amortized cost of key extraction is found to be about $\frac{2}{W+1}$ that of the Boneh-Franklin system, i.e., a threefold speedup for the common choice of 5-bit windows.

5 More Efficient Selective Identity IBE Based on BDHI Without Random Oracles

We construct an efficient IBE system that is selective identity, chosen plaintext secure without random oracles based on the Decision q -BDHI assumption (see Section 3.2). Decryption in the resulting IBE system is more efficient than the IBE construction in the previous section. Encryption efficiency and ciphertext size are the same.

5.1 Basic Construction

Let \mathbb{G} be a bilinear group of prime order p . For now, we assume that the public keys (ID) are elements in \mathbb{Z}_p^* . We show later that arbitrary identities in $\{0, 1\}^*$ can be used by first hashing ID using a collision resistant hash $H : \{0, 1\}^* \rightarrow \mathbb{Z}_p^*$. We also assume that the messages to be encrypted are elements in \mathbb{G}_1 . The IBE system works as follows:

Setup: To generate IBE parameters, select a random generator $g \in \mathbb{G}^*$, select random elements $x, y \in \mathbb{Z}_p^*$, and define $X = g^x$ and $Y = g^y$. The public parameters *params* and the secret *master-key* are given by

$$params = (g, X, Y), \quad master-key = (x, y)$$

KeyGen(*master-key*, ID): To create a private key for the public key $ID \in \mathbb{Z}_p^*$:

1. pick a random $r \in \mathbb{Z}_p$ and compute $K = g^{1/(ID+x+ry)} \in \mathbb{G}$,
2. output the private key $d_{ID} = (r, K)$.

In the unlikely event that $x + ry + ID = 0 \pmod{p}$, try again with a new random value for r .

Encrypt(*params*, ID, M): To encrypt a message $M \in \mathbb{G}_1$ under public key $ID \in \mathbb{Z}_p^*$, pick a random $s \in \mathbb{Z}_p^*$ and output the ciphertext

$$C = \left(g^{s \cdot ID} X^s, \quad Y^s, \quad e(g, g)^s \cdot M \right)$$

Note that $e(g, g)$ can be precomputed once and for all so that encryption does not require any pairing computations.

Decrypt(d_{ID} , C): To decrypt a ciphertext $C = (A, B, C)$ using the private key $d_{ID} = (r, K)$, output $C/e(AB^r, K)$. Indeed, for a valid ciphertext we have

$$\frac{C}{e(AB^r, K)} = \frac{C}{e(g^{s(ID+x+ry)}, g^{1/(ID+x+ry)})} = \frac{C}{e(g, g)^s} = M$$

Performance. In terms of efficiency, we note that the ciphertext size and encryption time are similar to the IBE system of the previous section. However, decryption requires only one pairing computation, as opposed to two in the previous section.

The IBE system above is related to a recent construction of Sakai and Kasahara [27, Sect. 3.1]. In the system of [27] the algorithm for generating user private keys is deterministic. In our system, key generation is randomized and this randomization is essential for the proof of security.

5.2 Proving Security

We prove security of the scheme under the Decision q -BDHI assumption from Section 3.2.

Theorem 5.1. *Suppose the (t, q, ϵ) -Decision BDHI assumption holds in \mathbb{G} of size $|\mathbb{G}| = p$. Then the previously defined IBE system is (t', q_s, ϵ) -selective identity, chosen plaintext (IND-sID-CPA) secure for any $q_s < q$, and any $t' < t - \Theta(\tau q^2)$ where τ is the maximum time for an exponentiation in \mathbb{G} .*

Proof. Suppose \mathcal{A} has advantage ϵ in attacking the IBE system. We build an algorithm \mathcal{B} that uses \mathcal{A} to solve the Decision q -BDHI problem in \mathbb{G} . Algorithm \mathcal{B} is given as input a random $(q+2)$ -tuple $(g, g^\alpha, g^{(\alpha^2)}, \dots, g^{(\alpha^q)}, T) \in (\mathbb{G}^*)^{q+1} \times \mathbb{G}_1$ that is either sampled from \mathcal{P}_{BDHI} (where $T = e(g, g)^{1/\alpha}$) or from \mathcal{R}_{BDHI} (where T is uniform and independent in \mathbb{G}_1). Algorithm \mathcal{B} 's goal is to output 1 if $T = e(g, g)^{1/\alpha}$ and 0 otherwise. Algorithm \mathcal{B} works by interacting with \mathcal{A} in a selective identity game as follows:

Preparation. Algorithm \mathcal{B} builds a generator $h \in \mathbb{G}^*$ for which it knows $q-1$ pairs of the form $(w_i, h^{1/(\alpha+w_i)})$ for random $w_1, \dots, w_{q-1} \in \mathbb{Z}_p^*$. This is done as follows:

1. Pick random $w_1, \dots, w_{q-1} \in \mathbb{Z}_p^*$ and let $f(z)$ be the polynomial $f(z) = \prod_{i=1}^{q-1} (z + w_i)$. Expand the terms of f to get $f(z) = \sum_{i=0}^{q-1} c_i z^i$. The constant term c_0 is non-zero.
2. Compute $h = \prod_{i=0}^{q-1} (g^{(\alpha^i)})^{c_i} = g^{f(\alpha)}$ and $u = \prod_{i=1}^q (g^{(\alpha^i)})^{c_{i-1}} = g^{\alpha f(\alpha)}$. Note that $u = h^\alpha$.
3. Check that $h \in \mathbb{G}^*$. Indeed if we had $h = 1$ in \mathbb{G} this would mean that $w_j = -\alpha$ for some easily identifiable w_j , at which point \mathcal{B} would be able to solve the challenge directly. We thus assume that all $w_j \neq -\alpha$.
4. Observe that for any $i = 1, \dots, q-1$, it is easy for \mathcal{B} to construct the pair $(w_i, h^{1/(\alpha+w_i)})$. To see this, write $f_i(z) = f(z)/(z + w_i) = \sum_{i=0}^{q-2} d_i z^i$. Then $h^{1/(\alpha+w_i)} = g^{f_i(\alpha)} = \prod_{i=0}^{q-2} (g^{(\alpha^i)})^{d_i}$.
5. Next, \mathcal{B} computes

$$T_h = T^{(c_0^2)} \cdot T_0 \quad \text{where} \quad T_0 = \prod_{i=0}^{q-1} \prod_{j=0}^{q-2} e\left(g^{(\alpha^i)}, g^{(\alpha^j)}\right)^{c_i c_{j+1}}$$

Observe that if $T = e(g, g)^{1/\alpha}$ then $T_h = e(g^{f(\alpha)/\alpha}, g^{f(\alpha)}) = e(h, h)^{1/\alpha}$. On the contrary, if T is uniform in \mathbb{G}_1 , then so is T_h .

We will be using the values h, u, T_h and the pairs $(w_i, h^{1/(\alpha+w_i)})$ for $i = 1, \dots, q-1$ throughout the simulation.

Initialization. The selective identity game begins with \mathcal{A} first outputting an identity $\text{ID}^* \in \mathbb{Z}_p^*$ that it intends to attack.

Setup. To generate the system parameters, algorithm \mathcal{B} does the following:

1. Pick random $a, b \in \mathbb{Z}_p^*$ under the constraint that $ab = \text{ID}^*$.
2. Compute $X = u^{-a}h^{-ab} = h^{-a(\alpha+b)}$ and $Y = u = h^\alpha$.
3. Publish $\text{params} = (h, X, Y)$ as the public parameters. Note that X, Y are independent of ID^* in the adversary's view.
4. We implicitly define $x = -a(\alpha + b)$ and $y = \alpha$ so that $X = h^x$ and $Y = h^y$. Algorithm \mathcal{B} does not know the value of x or y , but does know the value of $x + ay = -ab = -\text{ID}^*$.

Phase 1. \mathcal{A} issues up to $q_s < q$ private key queries. Consider the i -th query for the private key corresponding to public key $\text{ID}_i \neq \text{ID}^*$. We need to respond with a private key $(r, h^{1/(\text{ID}_i + x + ry)})$ for a uniformly distributed $r \in \mathbb{Z}_p$. Algorithm \mathcal{B} responds to the query as follows:

1. Let $(w_i, h^{1/(\alpha+w_i)})$ be the i -th pair constructed during the preparation step. Define $h_i = h^{1/(\alpha+w_i)}$.
2. \mathcal{B} first constructs an $r \in \mathbb{Z}_p$ satisfying $(r - a)(\alpha + w_i) = \text{ID}_i + x + ry$. Plugging in the values of x and y the equation becomes

$$(r - a)(\alpha + w_i) = \text{ID}_i - a(\alpha + b) + r\alpha$$

We see that the unknown α cancels from the equation and we get $r = a + \frac{\text{ID}_i - ab}{w_i} \in \mathbb{Z}_p$, which \mathcal{B} can evaluate.

3. Now, $(r, h_i^{1/(r-a)})$ is a valid private key for ID_i for two reasons. First,

$$h_i^{1/(r-a)} = (h^{1/(\alpha+w_i)})^{1/(r-a)} = h^{1/(r-a)(\alpha+w_i)} = h^{1/(\text{ID}_i + x + ry)}$$

as required. Second, r is uniformly distributed among all elements in \mathbb{Z}_p for which $\text{ID}_i + x + ry \neq 0$ and $r \neq a$. This is true since w_i is uniform in $\mathbb{Z}_p \setminus \{0, -\alpha\}$ and is currently independent of \mathcal{A} 's view. Algorithm \mathcal{B} gives \mathcal{A} the private key $(r, h_i^{1/(r-a)})$.

For completeness, we note that \mathcal{B} can construct the private key for ID_i with $r = a$ as $(r, h^{1/(\text{ID}_i - \text{ID}^*)})$. Hence, the r in the private key given to \mathcal{A} can be made uniform among all $r \in \mathbb{Z}_p$ for which $\text{ID}_i + x + ry \neq 0$ as required.

We point out that this procedure will fail to produce the private key for $\text{ID}_i = \text{ID}^*$ since in that case we get $r = a$ and $\text{ID}_i + x + ry = 0$. Hence, \mathcal{B} can generate private keys for *all* public keys except for ID^* .

Challenge. \mathcal{A} outputs two messages $M_0, M_1 \in \mathbb{G}_1$. Algorithm \mathcal{B} picks a random bit $b \in \{0, 1\}$ and a random $\ell \in \mathbb{Z}_p^*$. It responds with the ciphertext $CT = (h^{-a\ell}, h^\ell, T_h^\ell \cdot M_b)$. Define $s = \ell/\alpha$. On the one hand, if $T_h = e(h, h)^{1/\alpha}$ we have

$$\begin{aligned} h^{-a\ell} &= h^{-a\alpha(\ell/\alpha)} = h^{(x+ab)(\ell/\alpha)} = h^{(x+\text{ID}^*)(\ell/\alpha)} = h^{s\text{ID}^*} \cdot X^s \\ h^\ell &= Y^{\ell/\alpha} = Y^s \\ T_h^\ell &= e(h, h)^{\ell/\alpha} = e(h, h)^s \end{aligned}$$

It follows that CT is a valid encryption of M_b under ID^* , with the uniformly distributed randomization value $s = \ell/\alpha \in \mathbb{Z}_p^*$. On the other hand, when T_h is uniform in \mathbb{G}_1 , then, in the adversary's view, CT is independent of the bit b .

Phase 2. \mathcal{A} issues more private key queries, for a total of at most $q_s < q$. Algorithm \mathcal{B} responds as before.

Guess. Finally, \mathcal{A} outputs a guess $b' \in \{0, 1\}$. If $b = b'$ then \mathcal{B} outputs 1 meaning $T = e(g, g)^{1/\alpha}$. Otherwise, it outputs 0 meaning $T \neq e(g, g)^{1/\alpha}$.

We showed that when the input tuple is sampled from \mathcal{P}_{BDHI} (where $T = e(g, g)^{1/\alpha}$) then $T_h = e(h, h)^{1/\alpha}$ in which case \mathcal{A} must satisfy $|\Pr[b = b'] - 1/2| > \epsilon$. On the other hand, when the input tuple is sampled from \mathcal{R}_{BDHI} (where T is uniform in \mathbb{G}_1) then T_h is uniform and independent in \mathbb{G}_1 in which case $\Pr[b = b'] = 1/2$. Therefore, with g uniform in \mathbb{G}^* , x uniform in \mathbb{Z}_p^* , and T uniform in \mathbb{G}_1 , we have that

$$\left| \Pr \left[\mathcal{B}(g, g^x, \dots, g^{(x^q)}, e(g, g)^{1/x}) = 0 \right] - \Pr \left[\mathcal{B}(g, g^x, \dots, g^{(x^q)}, T) = 0 \right] \right| \geq \left| \left(\frac{1}{2} \pm \epsilon \right) - \frac{1}{2} \right| \geq \epsilon$$

as required. This completes the proof of Theorem 5.1. \square

5.3 Chosen Ciphertext Security

Canetti et al. [10, Section 2.2] describe a general method for converting a selective identity, chosen plaintext secure IBE into a selective identity, chosen ciphertext secure IBE. The method is based on [25, 26, 18]. Since it is generic, it applies to our system as well. In particular, the method can be used to render the IBE system above secure against chosen ciphertext attacks. The result is an IND-sID-CCA secure IBE without random oracles. However, the resulting system is inefficient since it relies on generic non-interactive zero-knowledge (NIZK) constructions.

5.4 Arbitrary Identities

As in the previous section, a standard argument shows that we can extend the IBE above to handle arbitrary identities $ID \in \{0, 1\}^*$ by first hashing ID using a collision resistant hash function $H : \{0, 1\}^* \rightarrow \mathbb{Z}_p^*$ prior to key generation and encryption. If the underlying scheme is selective identity, chosen plaintext (resp. ciphertext) secure, then so is the scheme with the additional hash function.

6 Efficient CCA2-Secure Public Key Systems

A recent result of Canetti et al. [11] gives a general method for constructing a CCA2 public key system from any selective identity, chosen plaintext IBE. Essentially the same result was used in Section 4 to transform our first HIBE construction into a chosen ciphertext secure HIBE of one lesser depth. The construction of [11] works by appending a one-time signature and a one-time signature public key to every ciphertext. Boneh and Katz [7] describe a more efficient transformation that requires only the addition of a MAC and a commitment to each ciphertext.

The [11, 7] transformations can be applied to the two IBE systems described in the previous two sections. In doing so, we obtain two new public key encryption schemes that are provably CCA2-secure without random oracles. We summarize here the performance characteristics of the two public key systems obtained from applying the [7] transformation.

- Encryption time: For both the Decision BDH system (Section 4) and the Decision BDHI system (Section 5) encryption time is dominated by three exponentiations in \mathbb{G} .

- Decryption time: For the Decision BDH system (Section 4) decryption time is dominated by the time to compute a product of two bilinear maps. For the Decision BDHI system (Section 5) decryption time is dominated by a single bilinear map computation. In both cases, the elements on the right side of all pairings do not depend on the ciphertext which enables further speed-up.
- Ciphertext size: For both systems the ciphertext is made up of three elements in \mathbb{G} plus a MAC and a commitment. The MAC and commitment together are about as long as one element in \mathbb{G} and hence total ciphertext size is about four elements in \mathbb{G} .

We note that encryption time is better than the original Cramer-Shoup [13] CCA2-secure public key which requires four exponentiations for encryption. Encryption time is the same as an improved version of Cramer-Shoup due to Kurosawa and Desmedt [17]. See [7] for a more detailed comparison.

A Non-Interactive CCA2-Secure Threshold Public Key System. We briefly note that in the IBE system of Section 4 it is easy to distribute the master key among n parties so that any t parties can be used to derive the private key for a given identity. When applying the technique of [11] to the resulting threshold IBE system, we obtain a CCA2-secure *threshold* public key system in the standard model: Given a ciphertext C , the combiner sends C to t of the n decryption parties. Each party checks the signature in C and, if it verifies, computes its share of the private key K needed to decrypt C (where K is an identity based private key in the underlying threshold IBE for a unique ID tied to the ciphertext). It sends the resulting share of the private key to the combiner who can then decrypt and recover the plaintext.

The resulting system is a CCA2-secure threshold public key system, without random oracles, in which there is no interaction needed between the decryption parties. Existing systems of this type, due to Canetti and Goldwasser [9], are based on the Cramer-Shoup system and require interaction between the decryption parties. The reason we are able to avoid interaction is that using the method of [11] anyone can check that a ciphertext is valid. In the Cramer-Shoup system only parties possessing the private key can check ciphertext validity, which makes threshold decryption non-trivial.

7 Fully Secure Identity Based Encryption

Until now we only discussed selective-ID security for IBE systems where the adversary commits ahead of time to the identity ID^* it wants to attack. In the full IBE security model [5] (denoted IND-ID-CPA) the attacker is allowed to *adaptively* choose which identity to attack by specifying ID^* in the challenge phase rather than in the setup phase. Giving the adversary more power this way makes it harder to construct fully secure IBE systems.

We briefly show that any selective-ID secure IBE is also a fully secure IBE, but the reduction is somewhat inefficient. First, we note that the selective-ID security of an IBE system is not weakened if additional restrictions on the identities are imposed (indeed, this only tightens the constraints on the adversary and relaxes those on the simulator). Identities in the systems of Sections 4 and 5 range natively over \mathbb{Z}_p and \mathbb{Z}_p^* , but by the preceding remark it is safe to restrict them to the set of integers $\{1, \dots, 2^n\}$ for $2^n < p$, represented as binary strings of length n . We can then expand our IBE schemes to arbitrary identities in $\{0, 1\}^*$ by first hashing identities using a collision resistant function with n -bit output, such as SHA-1 whose output is 160 bits. Hence, for an appropriately large p , taking $n = 160$ as the length of identities in the underlying IBE is a natural choice.

Let thus N be the number of allowed identities in the underlying IBE, where for example $N = 2^{160}$. The reduction from selective-ID IBE to fully secure IBE introduces a factor of N in

the security parameters of the system, as described in Theorem 7.1 below. Consequently, if the IBE system has sufficiently high selective-ID security (which requires using a bilinear group of sufficiently large size p) then the system is also a fully secure IBE with adequate security. This means that the selective-ID secure IBE system of [10] as well as the two systems described in the previous sections are fully secure IBE systems in their own right, assuming we use a large enough group so that the Decision BDH and Decision BDHI problems are sufficiently difficult. Extension to arbitrary identities requires collision resistant hashing, as mentioned in Sections 4.4 and 5.4, in which case N must be at least 2^{160} .

Concretely, an immediate corollary of Theorem 7.1 below is that, using 160-bits identities and using a group where no t -time adversary can break Decision BDH with advantage 2^{-240} , the IBE system of Section 4 is a $(t, q_S, 2^{-80})$ -fully secure IBE for any q_S . The system can be expanded to arbitrary identities in $\{0, 1\}^*$ by first hashing identities using a collision resistant hash function with a 160-bit output.

Theorem 7.1. *Let \mathcal{E} be a (t, q_S, ϵ) -selective identity secure IBE system (IND-sID-CPA). Suppose \mathcal{E} admits N distinct identities. Then \mathcal{E} is also a $(t, q_S, N\epsilon)$ -fully secure IBE (IND-ID-CPA).*

Proof. Suppose algorithm \mathcal{A} has advantage $N\epsilon$ in breaking the full security of the IBE system. We build an algorithm \mathcal{B} that has advantage ϵ in breaking selective-ID security of the system. Algorithm \mathcal{B} works as follows:

Init. \mathcal{B} picks a random $ID^* \in \{0, 1\}^n$ and outputs it as the identity that it wishes to attack.

Setup. The challenger gives \mathcal{B} the public parameters for an IBE system. \mathcal{B} forwards these parameters to \mathcal{A} .

Phase 1. \mathcal{A} issues private key queries. Consider the i 'th query for identity ID_i . If $ID_i \neq ID^*$ algorithm \mathcal{B} forwards the query to its challenger. Since the query is valid ($ID_i \neq ID^*$), the challenger responds with the private key for ID_i which \mathcal{B} then forwards to \mathcal{A} . However, in the unlikely event that $ID_i = ID^*$, algorithm \mathcal{B} cannot respond to this query. In this case, \mathcal{B} terminates the simulation, picks a random bit b' , and outputs b' as its guess for the challenger's bit b .

Challenge. Once phase 1 is over \mathcal{A} outputs an identity $ID_0^* \in \{0, 1\}^n$ and two equal length messages M_0, M_1 . Algorithm \mathcal{B} forwards M_0, M_1 to its challenger and receives back the challenge ciphertext C^* . We consider two cases:

1. If $ID^* \neq ID_0^*$ then algorithm \mathcal{B} picks a random bit $b' \in \{0, 1\}$, outputs b' as its guess for b , and terminates.
2. Otherwise, $ID^* = ID_0^*$ in which case C^* is a proper encryption of one of M_0 or M_1 under ID_0^* as expected by \mathcal{A} . Algorithm \mathcal{B} gives C^* to \mathcal{A} and continues to Phase 2.

Phase 2. \mathcal{A} continues to issue private key queries. \mathcal{B} responds as before. Since now the queries cannot equal ID^* these queries cannot cause \mathcal{B} to abort.

Output. Finally, \mathcal{A} outputs its guess $b' \in \{0, 1\}$ for b . \mathcal{B} outputs the same b' as its guess for b .

Next, we analyze \mathcal{B} 's advantage in guessing b . Let $q_1 \leq q_S < N$ be the number of distinct queries that \mathcal{A} issued during phase 1. Let success_1 denote the event that during phase 1 \mathcal{A} did not issue a query for ID^* . Let success denote the event that both success_1 occurred and $ID^* = ID_0^*$. Then

$$\Pr[\text{success}] = \Pr[\text{success}_1] \cdot \Pr[ID^* = ID_0^* \mid \text{success}_1] = \left(1 - \frac{q_1}{N}\right) \frac{1}{N - q_1} = \frac{1}{N}$$

When event `success` happens, \mathcal{A} 's view is identical to its view in a real attack game and therefore $|\Pr[b = b'|\text{success}] - \frac{1}{2}| \geq N\epsilon$. Furthermore, by definition of \mathcal{B} we have that $\Pr[b = b'|\overline{\text{success}}] = \frac{1}{2}$. It follows that

$$\begin{aligned} \left| \Pr[b = b'] - \frac{1}{2} \right| &= \left| \left(\Pr[b = b'|\text{success}] \cdot \Pr[\text{success}] \right) + \left(\Pr[b = b'|\overline{\text{success}}] \cdot \Pr[\overline{\text{success}}] \right) - \frac{1}{2} \right| \\ &= \left| \Pr[b = b'|\text{success}] \cdot \frac{1}{N} + \frac{1}{2} \cdot \frac{N-1}{N} - \frac{1}{2} \right| = \left| \Pr[b = b'|\text{success}] - \frac{1}{2} \right| \cdot \frac{1}{N} \geq \epsilon \end{aligned}$$

as required. This completes the proof of Theorem 7.1. \square

A natural question is whether one can build a fully secure IBE with a more efficient security reduction than in Theorem 7.1. Building on the system of Section 4 we were recently able to construct a fully secure IBE where the security reduction only introduces an error factor of $\tilde{O}(q_s^2)$, as oppsed to N , as described in [4]. However the construction is not very practical and mostly serves as a proof of concept. Very recently, Waters [32] was able to construct a very elegant generalization of the system of Section 4 that gives an efficient fully secure IBE where the security reduction introduces an error factor of roughly $O(q_s)$.

Fully Secure IBE Using Random Oracles. It is also worth noting that a random oracle H immediately converts a selective-ID IBE \mathcal{E} to a fully secure IBE by the process of hashing the identity ID with H before using ID . We denote the resulting system by \mathcal{E}_H . We state this in the following theorem.

Theorem 7.2. *Let \mathcal{E} be a (t, q_s, ϵ) selective-ID secure IBE. Suppose identities in \mathcal{E} are n -bits long. Let H be a hash function $H : \{0, 1\}^* \rightarrow \{0, 1\}^n$ modeled as a random oracle. Then \mathcal{E}_H is a (t, q_s, ϵ') fully secure IBE (in the random oracle model) for $\epsilon' = \epsilon \cdot q_H / (1 - q_s/2^n) \approx q_H \cdot \epsilon$, where q_H is the maximum number of oracle calls to H that the adversary can make.*

The proof of the theorem is similar to the proof of Theorem 7.1 and is omitted. We merely mention that the additional correction factor $(1 - q_s/2^n)^{-1}$ accounts for the possible random oracle collisions. It is worth noting that in the proof, the random oracle is “programmed” at only one point. Unlike the original Boneh-Franklin IBE scheme and its many variants, where the random oracle programmability is crucial to answer all private key queries, here the random oracle is only needed to ensure that the hash of the challenge identity provided by the adversary (ID_0^*) is mapped to the identity chosen by the simulator at the initialization step (ID^*).

A consequence of Theorems 7.1 and 7.2 is that, using a collision resistant function $H : \{0, 1\}^* \rightarrow \{0, 1\}^n$ for a properly chosen n satisfying $2^{160} \leq 2^n \ll p$, the same hash IBE scheme \mathcal{E}_H can be proven secure in both the standard model and the random oracle model. Specifically, the scheme features $(t, q_s, 2^n\epsilon)$ -full IBE security in the standard model, which is boosted to $(t, q_s, q_H\epsilon)$ when H is viewed as a random oracle.

8 Arbitrary Length Identities Without Collision Resistance

We mentioned in Sections 4.4 and 5.4 that both schemes can easily be extended to accept arbitrary length strings as public keys, using collision resistant hashing. It is well known that collision resistant hash functions can be constructed in groups where the discrete logarithm problem is hard; which is clearly the case under the present assumptions. However, in practice, it is advantageous

to use faster hash constructions such as SHA-1, whose collision resistance rests mainly on faith. In this section we ask whether our IBE systems can be extended to accept arbitrary length identities using a weaker assumption than collision resistance. We observe that a universal one-way hash family (UOWHF) of functions $H_k : \{0, 1\}^* \rightarrow \mathcal{I}$ is sufficient to transform any selective-ID secure IBE with a finite identity set \mathcal{I} into a selective-ID secure IBE that takes arbitrary length identities in $\{0, 1\}^*$.

Recall that a (t', ϵ') -UOWHF [24] is a family of functions $H_k : \{0, 1\}^* \rightarrow \mathcal{I}$ with indices in \mathcal{K} , such that, for any randomized algorithm $(\mathcal{A}_1, \mathcal{A}_2)$ running in time t' , it holds that:

$$\Pr [H_k(x) = H_k(x') : (x, \sigma) \leftarrow \mathcal{A}_1, k \xleftarrow{R} \mathcal{K}, x' \leftarrow \mathcal{A}_2(x, k, \sigma)] < \epsilon'$$

The probability is taken over the random choice of k and the random coins used by $(\mathcal{A}_1, \mathcal{A}_2)$. For practical purposes, it is assumed that the index k ranges in a finite set \mathcal{K} and has a compact representation.

Let \mathcal{E} be an IBE system that admits identities $ID \in \mathcal{I}$. Suppose that $\mathcal{H} = \{H_k\}$ is a (t, ϵ') -UOWHF of functions $H_k : \{0, 1\}^* \rightarrow \mathcal{I}$ indexed by $k \in \mathcal{K}$. We construct a new IBE system \mathcal{E}' that takes arbitrary length identities $ID \in \{0, 1\}^*$ as follows:

Setup $_{\mathcal{E}'}$ runs **Setup $_{\mathcal{E}}$** to obtain *params* and *master-key*, and appends a random $k \in \mathcal{K}$ to give *params'* = (*params*, k). The new public *params'* and the original *master-key* are returned.

KeyGen $_{\mathcal{E}'}$ (*params'*, *master-key*, ID') returns $d_{ID'} \leftarrow \text{KeyGen}_{\mathcal{E}}(\text{params}, H_k(ID'))$.

Encrypt $_{\mathcal{E}'}$ (*params'*, ID' , M) returns $\text{Encrypt}_{\mathcal{E}}(\text{params}, ID, M)$ where $ID = H_k(ID')$.

Decrypt $_{\mathcal{E}'}$ (*params'*, $d_{ID'}$, C) returns $\text{Decrypt}_{\mathcal{E}}(\text{params}, d_{ID'}, C)$.

The following theorem shows that \mathcal{E}' is a selective identity secure IBE system. The proof is immediate and is omitted.

Theorem 8.1. *Let \mathcal{E} be a (t, q_S, ϵ) -selective identity secure IBE system (IND-sID-CPA) that admits identities $ID \in \mathcal{I}$. Suppose that $\mathcal{H} = \{H_k\}$ is a (t, ϵ') -UOWHF of functions $H_k : \{0, 1\}^* \rightarrow \mathcal{I}$ indexed by $k \in \mathcal{K}$. Then, \mathcal{E}' is a $(t, q_S, \epsilon + \epsilon')$ -selective identity secure IBE system that admits identities $ID' \in \{0, 1\}^*$.*

9 DHI and Generalized Diffie-Hellman

In Section 3.2 we defined the q -BDHI problem in a bilinear group. A closely related problem is the q -Diffie-Hellman Inversion (q -DHI) problem: given a tuple $(g, g^x, g^{(x^2)}, \dots, g^{(x^q)}) \in \mathbb{G}^{q+1}$ as input, output $g^{1/x} \in \mathbb{G}$. Here, \mathbb{G} need not be a bilinear group. Loosely speaking, the q -DHI assumption states that the q -DHI problem is intractable in \mathbb{G} . This assumption was previously used in [22] where it was called weak Diffie-Hellman.

Many cryptographic constructions rely on the Generalized Diffie-Hellman assumption (GenDH) for security [21, 23, 2, 19, 8]. In this section we show that the q -DHI assumption implies the $(q+1)$ -Generalized Diffie-Hellman assumption. Thus, constructions that rely on Generalized Diffie-Hellman could instead rely on q -DHI which appears to be a more natural complexity assumption, and is easier to state since the problem description does not require an oracle.

We first review the GenDH assumption. The assumption says that, for a random generator g of \mathbb{G} , given g^{a_1}, \dots, g^{a_q} in \mathbb{G} and given all the subset products $g^{\prod_{i \in S} a_i} \in \mathbb{G}$ for any strict

subset $S \subset \{1, \dots, q\}$, it is hard to compute $g^{a_1 \dots a_q} \in \mathbb{G}$. Since the number of subset products is exponential in q , access to all these subset products is provided through an oracle. For a vector $\vec{a} = (a_1, \dots, a_q) \in \mathbb{Z}_p^q$, define $\mathcal{O}_{g, \vec{a}}$ to be an oracle that for any strict subset $S \subset \{1, \dots, q\}$ responds with

$$\mathcal{O}_{g, \vec{a}}(S) = g^{\prod_{i \in S} a_i} \in \mathbb{G}.$$

Define the advantage of algorithm \mathcal{A} in solving the generalized Diffie-Hellman problem to be the probability that \mathcal{A} is able to compute $g^{a_1 \dots a_q}$ given access to the oracle $\mathcal{O}_{g, \vec{a}}(S)$. In other words,

$$\text{Adv}_{\mathcal{A}, q} = \Pr[\mathcal{A}^{\mathcal{O}_{g, \vec{a}}} = g^{a_1 \dots a_q} : g \leftarrow \mathbb{G}^*, \vec{a} = (a_1, \dots, a_q) \leftarrow \mathbb{Z}_p^q]$$

Note that the oracle only answers queries for strict subsets of $\{1, \dots, q\}$.

Definition 9.1. We say that \mathbb{G} satisfies the (t, q, ϵ) -Generalized Diffie-Hellman assumption if for all t -time algorithms \mathcal{A} we have $\text{Adv}_{\mathcal{A}, q} < \epsilon$.

Theorem 9.2. Suppose the $(t, q-1, \epsilon)$ -DHI assumption holds in \mathbb{G} . Then the (t, q, ϵ) -GenDH assumption also holds in \mathbb{G} .

Proof. Suppose \mathcal{A} is an algorithm that has advantage ϵ in solving the q -GenDH problem. We construct an algorithm \mathcal{B} that solves $(q-1)$ -DHI with the same advantage ϵ . Algorithm \mathcal{B} is given $g, g^x, g^{(x^2)}, \dots, g^{(x^{q-1})} \in \mathbb{G}$ and its goal is to compute $g^{1/x} \in \mathbb{G}$. Let $h = g^{(x^{q-1})}$ and $y = x^{-1} \in \mathbb{Z}_p$. Then the input to \mathcal{B} can be re-written as $h, h^y, h^{(y^2)}, \dots, h^{(y^{q-1})} \in \mathbb{G}$ and \mathcal{B} 's goal is to output $h^{(y^q)} = g^{1/x}$.

Algorithm \mathcal{B} first picks q random values $c_1, \dots, c_q \in \mathbb{Z}_p$. It then runs algorithm \mathcal{A} and simulates the oracle $\mathcal{O}_{h, \vec{a}}$ for \mathcal{A} . The vector \vec{a} that \mathcal{B} will use is $\vec{a} = (y + c_1, \dots, y + c_q)$. Note that \mathcal{B} does not know \vec{a} explicitly since \mathcal{B} does not have y . When \mathcal{A} issues a query for $\mathcal{O}_{h, \vec{a}}(S)$ for some strict subset $S \subset \{1, \dots, q\}$ algorithm \mathcal{B} responds as follows:

1. Define the polynomial $f(z) = \prod_{i \in S} (z + c_i)$ and expand the terms to obtain $f(z) = \sum_{i=0}^{|S|} b_i z^i$.
2. Compute $t = \prod_{i=0}^{|S|} (h^{(y^i)})^{b_i} = h^{f(y)}$. Since $|S| < q$ all the values $h^{(y^i)}$ in the product are known to \mathcal{B} .
3. By construction we know that $t = h^{\prod_{i \in S} (y + c_i)}$. Algorithm \mathcal{B} responds by setting $\mathcal{O}_{h, \vec{a}}(S) = t$.

The responses to all of the adversary's oracle queries are consistent with the hidden vector $\vec{a} = (y + c_1, \dots, y + c_q)$. Therefore, eventually, \mathcal{A} will output $T = h^{\prod_{i=1}^q (y + c_i)}$. Define the polynomial $f(z) = \prod_{i=1}^q (z + c_i)$ and expand the terms to get $f(z) = z^q + \sum_{i=0}^{q-1} b_i z^i$. To conclude, \mathcal{B} outputs

$$T / \prod_{i=0}^{q-1} (h^{(y^i)})^{b_i} = h^{(y^q)}$$

which is the required value. □

The same property as in Theorem 9.2 also holds for the decision versions of the DHI and GenDH problems. The q -DHI assumption is easier to state than the q -GenDH assumption since there is no need for an oracle. When appropriate, constructions that depend on GenDH for security could instead use the DHI assumption.

10 Conclusions

We constructed two IBE systems that are secure against selective identity attacks in the standard model, i.e., without using random oracles. The first construction is based on the now classic BDH assumption. It extends readily to give a selective identity HIBE without random oracles, that can efficiently be made chosen ciphertext secure using a technique of [11]. The second construction is based on the Bilinear Diffie-Hellman Inversion assumption. The same technique of [11, 7] converts both our constructions into efficient CCA2-secure public key systems without random oracles that are almost as efficient as the Cramer-Shoup public key system.

We observed that a selective-ID secure IBE system implies a fully secure IBE system. Although the resulting security reduction is not polynomial, the parameters can easily be dimensioned in such a way as to result in a fully secure practical IBE system in the standard model, that is computationally as efficient, within a constant factor, as its random oracle counterpart of equivalent security.

Using the first construction in this paper as a building block, we were recently able to construct the first fully secure IBE system without the random oracle model with an efficient security reduction [4], thereby resolving an open problem posed in [5]. However, that system is quite impractical and should only be viewed as a constructive proof that such constructions are indeed possible. Waters [32] recently described an elegant generalization of this building block that directly gives a far more efficient fully secure IBE without the random oracle model. The question of constructing a fully secure IBE system with a tight reduction in the standard model remains open.

Acknowledgments

The comment on threshold systems in Section 6 came out of a discussion with Shai Halevi. We also thank Shai Halevi and Jonathan Katz for helpful comments on this work.

References

- [1] M. Bellare and P. Rogaway. Random oracle are practical: A paradigm for designing efficient protocols. In *Proceedings of the First ACM Conference on Computer and Communications Security*, pages 62–73, 1993.
- [2] E. Biham, D. Boneh, and O. Reingold. Breaking generalized Diffie-Hellman modulo a composite is no easier than factoring. *Information Processing Letters*, 70:83–87, 1999.
- [3] D. Boneh and X. Boyen. Efficient selective-ID identity based encryption without random oracles. In *Advances in Cryptology—EUROCRYPT 2004*, volume 3027 of *LNCS*, pages 223–38. Springer-Verlag, 2004.
- [4] D. Boneh and X. Boyen. Secure identity based encryption without random oracles. In M. Franklin, editor, *Advances in Cryptology—CRYPTO 2004*, volume 3152 of *LNCS*, pages 443–59. Springer-Verlag, 2004.
- [5] D. Boneh and M. Franklin. Identity-based encryption from the Weil pairing. In J. Kilian, editor, *Advances in Cryptology—CRYPTO 2001*, volume 2139 of *LNCS*, pages 213–29. Springer-Verlag, 2001.

- [6] D. Boneh and M. Franklin. Identity-based encryption from the Weil pairing. *SIAM Journal of Computing*, 32(3):586–615, 2003.
- [7] D. Boneh and J. Katz. Improved efficiency for CCA-secure cryptosystems built using identity based encryption. Submitted for publication, 2004.
- [8] D. Boneh and A. Silverberg. Applications of multilinear forms to cryptography. *Contemporary Mathematics*, 324:71–90, 2003.
- [9] R. Canetti and S. Goldwasser. An efficient threshold public key cryptosystem secure against adaptive chosen ciphertext attack. In *Advances in Cryptology—EUROCRYPT 1999*, pages 90–106, 1999.
- [10] R. Canetti, S. Halevi, and J. Katz. A forward-secure public-key encryption scheme. In *Advances in Cryptology—EUROCRYPT 2003*, volume 2656 of *LNCS*. Springer-Verlag, 2003.
- [11] R. Canetti, S. Halevi, and J. Katz. Chosen-ciphertext security from identity-based encryption. In *Advances in Cryptology—EUROCRYPT 2004*, volume 3027 of *LNCS*, pages 207–22. Springer-Verlag, 2004.
- [12] C. Cocks. An identity based encryption scheme based on quadratic residues. In *Proceedings of the 8th IMA International Conference on Cryptography and Coding*, pages 26–8, 2001.
- [13] R. Cramer and V. Shoup. A practical public key cryptosystem provably secure against adaptive chosen ciphertext attacks. In H. Krawczyk, editor, *Proceedings of Crypto 1998*, volume 1462 of *LNCS*, pages 13–25. Springer-Verlag, 1998.
- [14] C. Gentry and A. Silverberg. Hierarchical ID-based cryptography. In *Advances in Cryptology—ASIACRYPT 2002*, 2002.
- [15] J. Horwitz and B. Lynn. Towards hierarchical identity-based encryption. In *Advances in Cryptology—EUROCRYPT 2002*, pages 466–81, 2002.
- [16] A. Joux. A one round protocol for tripartite Diffie-Hellman. In W. Bosma, editor, *Proceedings of ANTS IV*, volume 1838 of *LNCS*, pages 385–94. Springer-Verlag, 2000.
- [17] K. Kurosawa and Y. Desmedt. A new paradigm of hybrid encryption scheme. In *Advances in Cryptology—CRYPTO 2004*, volume 3152 of *LNCS*, pages 426–42. Springer-Verlag, 2004.
- [18] Y. Lindell. A simpler construction of CCA2-secure public-key encryption under general assumptions. In *Advances in Cryptology—EUROCRYPT 2003*, volume 2656 of *LNCS*, pages 241–54, 2003.
- [19] A. Lysyanskaya. Unique signatures and verifiable random functions from the DH-DDH separation. In *Advances in Cryptology—CRYPTO 2002*, LNCS. Springer-Verlag, 2002.
- [20] U. M. Maurer and Y. Yacobi. A non-interactive public-key distribution system. *Designs, Codes and Cryptography*, 9(3):305–316, November 1996.
- [21] G. T. Michael Steiner and M. Waidner. Diffie-Hellman key distribution extended to groups. In *Proceedings of the ACM Conference on Computer and Communications Security*, 1996.

- [22] S. Mitsunari, R. Sakai, and M. Kasahara. A new traitor tracing. *IEICE Transactions on Fundamentals*, E85-A(2):481–4, 2002.
- [23] M. Naor and O. Reingold. Number-theoretic constructions of efficient pseudo-random functions. In *Proceedings of the 38th IEEE Symposium on Foundations of Computer Science*, pages 458–67, 1997.
- [24] M. Naor and M. Yung. Universal one-way hash functions and their cryptographic applications. In *Proceedings of the 21st ACM Symposium on Theory of Computing*, 1989.
- [25] M. Naor and M. Yung. Public key cryptosystems provable secure against chosen ciphertext attacks. In *Proceedings of the ACM Symposium on Theory of Computing*, pages 427–37. ACM, 1990.
- [26] A. Sahai. Non-malleable non-interactive zero knowledge and adaptive chosen-ciphertext security. In *Proceedings of the 40th IEEE Symposium on Foundations of Computer Science*, 1999.
- [27] R. Sakai and M. Kasahara. ID based cryptosystems with pairing over elliptic curve, 2003. <http://eprint.iacr.org/2003/054>.
- [28] R. Sakai, K. Ohgishi, and M. Kasahara. Cryptosystems based on pairings. In *Proceedings of the Symposium on Cryptography and Information Security—SCIS 2000*, Japan, 2000.
- [29] A. Shamir. Identity-based cryptosystems and signature schemes. In *Advances in Cryptology—CRYPTO 1984*, volume 196 of *LNCS*, pages 47–53. Springer-Verlag, 1984.
- [30] H. Tanaka. A realization scheme for the identity-based cryptosystem. In *Advances in Cryptology—CRYPTO 1987*, volume 293 of *LNCS*, pages 341–349. Springer-Verlag, 1987.
- [31] S. Tsujii and T. Itoh. An ID-based cryptosystem based on the discrete logarithm problem. *IEEE Journal on Selected Areas in Communication*, 7(4):467–473, 1989.
- [32] B. Waters. Efficient identity-based encryption without random oracles. Preprint, 2004. <http://eprint.iacr.org/2004/180>.